
How Many Clusters? Which Clustering Method? Answers Via Model-Based Cluster Analysis

CHRIS FRALEY AND ADRIAN E. RAFTERY

Department of Statistics, University of Washington, USA
Email: fraley@stat.washington.edu

We consider the problem of **determining the structure of clustered data, without prior knowledge of the number of clusters or any other information about their composition**. Data are represented by a mixture model in which each component corresponds to a different cluster. Models with varying geometric properties are obtained through Gaussian components with different parametrizations and cross-cluster constraints. Noise and outliers can be modelled by adding a Poisson process component. Partitions are determined by the expectation-maximization (EM) algorithm for maximum likelihood, with initial values from agglomerative hierarchical clustering. Models are compared using an approximation to the Bayes factor based on the **Bayesian information criterion** (BIC); unlike significance tests, this allows comparison of more than two models at the same time, and removes the restriction that the models compared be nested. The problems of determining the number of clusters and the clustering method are solved simultaneously by choosing the best model. Moreover, the EM result provides a measure of uncertainty about the associated classification of each data point. Examples are given, showing that this approach can give performance that is much better than standard procedures, which often fail to identify groups that are either overlapping or of varying sizes and shapes.

1. INTRODUCTION

We consider the problem of determining the intrinsic structure of clustered data when no information other than the observed values is available. This problem is known as *cluster analysis* and should be distinguished from the related problem of *discriminant analysis*, in which known groupings of some observations are used to categorize others and infer the structure of the data as a whole.

Probability models have been proposed for quite some time as a basis for cluster analysis. In this approach, the data are viewed as coming from a mixture of probability distributions, each representing a different cluster. Recently, methods of this type have shown promise in a number of practical applications, including character recognition (Murtagh and Raftery [1]), tissue segmentation (Banfield and Raftery [2]), minefield and seismic fault detection (Dasgupta and Raftery [3]), identification of textile flaws from images (Campbell *et al.* [4]) and classification of astronomical data (Celeux and Govaert [5], Mukerjee *et al.* [6]).

Bayes factors, approximated by the **Bayesian information criterion (BIC)**, have been applied successfully to the problem of determining the number of components in a model [3, 6] and for deciding which among two or more partitions most closely matches the data for a given model [4]. We describe a clustering methodology based on multivariate normal mixtures in which the BIC is used

for direct comparison of models that may differ not only in the number of components in the mixture, but also in the underlying densities of the various components. Partitions are determined (as in [3]) by a combination of hierarchical clustering and the expectation-maximization (EM) algorithm (Dempster, Laird and Rubin [7]) for maximum likelihood. This approach can give much better performance than existing methods. Moreover, the EM result also provides a measure of uncertainty about the resulting classification. Figure 1 shows an example in which model-based classification is able to match the clinical classification of a biomedical data set much more closely than single-link (nearest-neighbour) or standard k -means, in the absence of any training data.

This paper is organized as follows. In Section 2, we give the necessary background in multivariate cluster analysis, including discussions of probability models, the EM algorithm for clustering and approximate Bayes factors. The basic model-based strategy and modifications for handling noise are described in Sections 2.5 and 2.6, respectively. A detailed analysis of the multivariate data set shown in Figure 1 is given in Section 3.1, followed by an example from minefield detection in the presence of noise in Section 3.2. Information on available software for the various procedures used in this approach is given in Section 4. A final section summarizes and indicates extensions to the method.

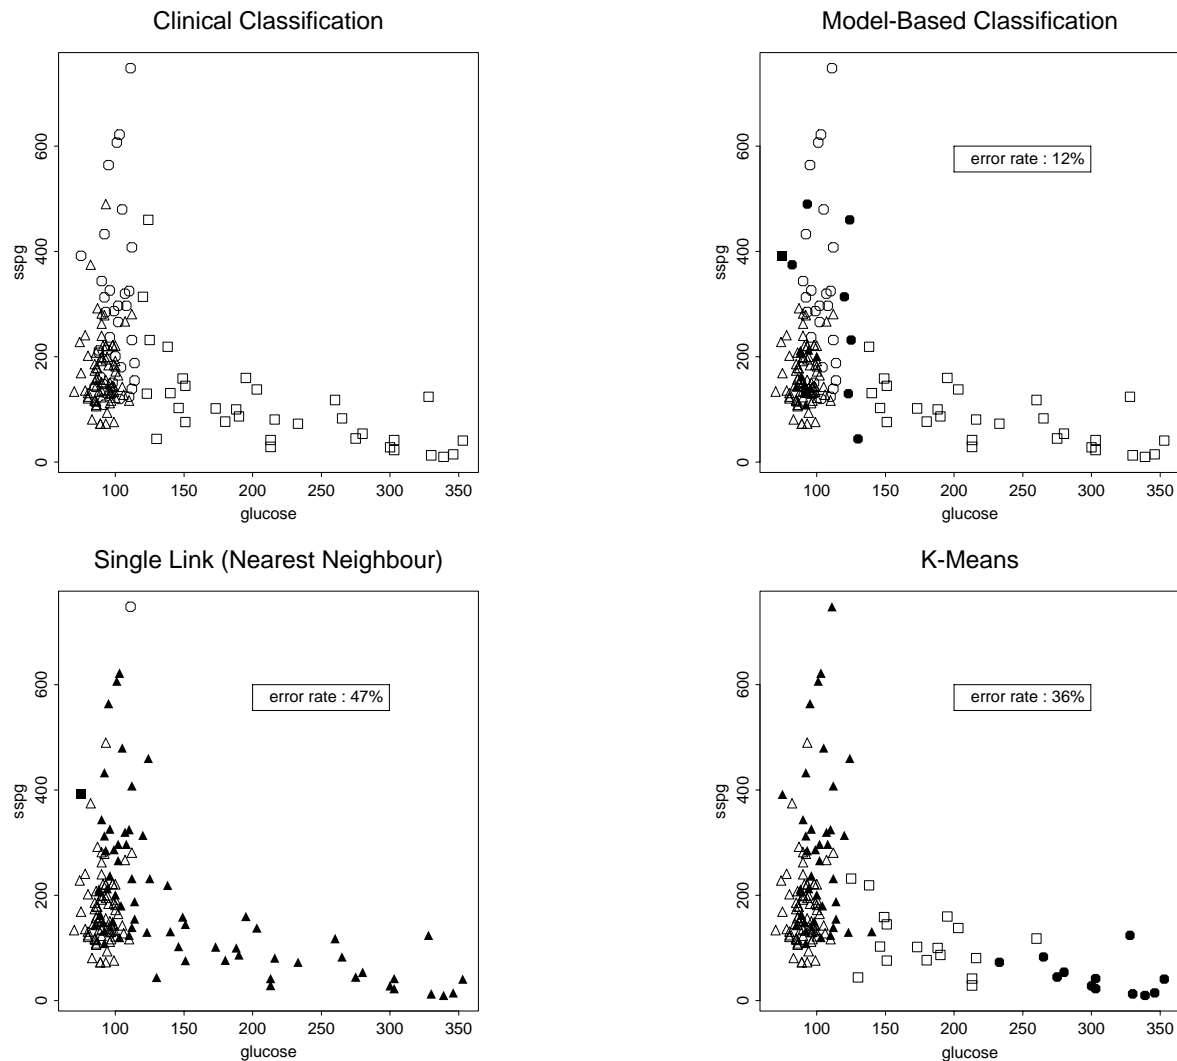


FIGURE 1. A projection of the three-group classification of the diabetes data from Reaven and Miller [8] using single link or nearest neighbour, standard k -means and the unconstrained model-based approach. Filled symbols represent misclassified observations.

2. MODEL-BASED CLUSTER ANALYSIS

2.1. Cluster analysis background

By *cluster analysis* we mean the partitioning of data into meaningful subgroups, when the number of subgroups and other information about their composition may be unknown; good introductions include Hartigan [9], Gordon [10], Murtagh [11], McLachlan and Basford [12] and Kaufman and Rousseeuw [13]. Clustering methods range from those that are largely heuristic to more formal procedures based on statistical models. They usually follow either a hierarchical strategy or one in which observations are relocated among tentative clusters.

Hierarchical methods proceed by stages producing a sequence of partitions, each corresponding to a different number of clusters. They can be either 'agglomerative', meaning that groups are merged, or 'divisive', in which one or more groups are split at each stage. Hierarchical procedures that use subdivision are not practical unless the number of possible splittings can somehow be restricted.

In agglomerative hierarchical clustering, however, the number of stages is bounded by the number of groups in the initial partition. It is common practice to begin with each observation in a cluster by itself, although the procedure could be initialized from a coarser partition if some groupings are known. A drawback of agglomerative methods is that those that are practical in terms of time efficiency require memory usage proportional to the square of the number of groups in the initial partition.

At each stage of hierarchical clustering, the splitting or merging is chosen so as to optimize some criterion. Conventional agglomerative hierarchical methods use heuristic criteria, such as single link (nearest neighbour), complete link (farthest neighbour) or sum of squares [13]. In model-based methods, a maximum-likelihood criterion is used for merging groups [1, 2].

Relocation methods move observations iteratively from one group to another, starting from an initial partition. The number of groups has to be specified in advance and typically does not change during the course of the

iteration. The most common relocation method— k -means (MacQueen [14], Hartigan and Wong [15])—reduces the within-group sums of squares. For clustering via mixture models, relocation techniques are usually based on the EM algorithm [7] (see Section 2.3).

Neither hierarchical nor relocation methods directly address the issue of determining the number of groups within the data. Various strategies for simultaneous determination of the number of clusters and cluster membership have been proposed (e.g. Engelman and Hartigan [16], Bock [17], Bozdogan [18]—for a survey see Bock [19]). An alternative is described in this paper.

2.2. Probability models for cluster analysis

In model-based clustering, it is assumed that the data are generated by a mixture of underlying probability distributions in which each component represents a different group or cluster. Given observations $\mathbf{x} = (\mathbf{x}_1, \dots, \mathbf{x}_n)$, let $f_k(\mathbf{x}_i | \theta_k)$ be the density of an observation \mathbf{x}_i from the k th component, where θ_k are the corresponding parameters, and let G be the number of components in the mixture. The model for the composite of the clusters is usually formulated in one of two ways. The *classification likelihood* approach maximizes

$$\mathcal{L}_C(\theta_1, \dots, \theta_G; \gamma_1, \dots, \gamma_n | \mathbf{x}) = \prod_{i=1}^n f_{\gamma_i}(\mathbf{x}_i | \theta_{\gamma_i}), \quad (1)$$

where γ_i are discrete values labelling the classification: $\gamma_i = k$ if \mathbf{x}_i belongs to the k th component. The *mixture likelihood* approach maximizes

$$\mathcal{L}_M(\theta_1, \dots, \theta_G; \tau_1, \dots, \tau_G | \mathbf{x}) = \prod_{i=1}^n \sum_{k=1}^G \tau_k f_k(\mathbf{x}_i | \theta_k), \quad (2)$$

where τ_k is the probability that an observation belongs to the k th component ($\tau_k \geq 0$; $\sum_{k=1}^G \tau_k = 1$).

We are mainly concerned with the case where $f_k(\mathbf{x}_i | \theta_k)$ is multivariate normal (Gaussian), a model that has been used with considerable success in a number of applications [1, 2, 5, 3, 4, 6]. In this instance, the parameters θ_k consist of a mean vector μ_k and a covariance matrix Σ_k , and the density has the form

$$f_k(\mathbf{x}_i | \mu_k, \Sigma_k) = \frac{\exp\{-\frac{1}{2}(\mathbf{x}_i - \mu_k)^T \Sigma_k^{-1}(\mathbf{x}_i - \mu_k)\}}{(2\pi)^{p/2} |\Sigma_k|^{1/2}}. \quad (3)$$

Clusters are ellipsoidal, centred at means μ_k . The covariances Σ_k determine their other geometric characteristics.

Banfield and Raftery [2] developed a model-based framework for clustering by parametrizing the covariance matrix in terms of its eigenvalue decomposition in the form

$$\Sigma_k = \lambda_k D_k A_k D_k^T, \quad (4)$$

where D_k is the orthogonal matrix of eigenvectors, A_k is a diagonal matrix whose elements are proportional to the

eigenvalues of Σ_k and λ_k is a scalar. The orientation of the principal components of Σ_k is determined by D_k , while A_k determines the shape of the density contours; λ_k specifies the volume of the corresponding ellipsoid, which is proportional to $\lambda_k^p |A_k|$.¹ Characteristics (orientation, volume and shape) of distributions are usually estimated from the data and can be allowed to vary between clusters or constrained to be the same for all clusters.

This approach subsumes several earlier proposals based on Gaussian mixtures: $\Sigma_k = \lambda I$ gives the sum of squares criterion, long known as a heuristic (Ward [20]), in which clusters are spherical and have equal volumes; $\Sigma_k = \Sigma = \lambda D A D^T$, in which all clusters have the same shape, volume and orientation (Friedman and Rubin [21]); unconstrained $\Sigma_k = \lambda_k D_k A_k D_k^T$, which is the most general model (Scott and Symons [22]); and $\Sigma_k = \lambda D_k A D_k$ (Murtagh and Raftery [1]), in which only the orientations of the clusters may differ. Table 1 shows the geometric interpretation of the various parametrizations discussed in [2]. A more extensive set of models within the same framework is treated in [5].

The classification likelihood can be used as the basis for agglomerative hierarchical clustering [1, 2]. At each stage, a pair of clusters is merged so as to maximize the resulting likelihood. Fraley [23] developed efficient algorithms for hierarchical clustering with the various parametrizations (4) of Gaussian mixture models.

2.3. EM algorithms for clustering

Iterative relocation methods for clustering via mixture models are possible through EM and related techniques [12]. The EM algorithm [7, 24] is a general approach to maximum likelihood in the presence of incomplete data. In EM for clustering, the ‘complete’ data are considered to be $\mathbf{y}_i = (\mathbf{x}_i, \mathbf{z}_i)$, where $\mathbf{z}_i = (z_{i1}, \dots, z_{iG})$ with

$$z_{ik} = \begin{cases} 1 & \text{if } \mathbf{x}_i \text{ belongs to group } k \\ 0 & \text{otherwise} \end{cases} \quad (5)$$

constitutes the ‘missing’ data. The relevant assumptions are that the density of an observation \mathbf{x}_i given \mathbf{z}_i is given by $\prod_{k=1}^G f_k(\mathbf{x}_i | \theta_k)^{z_{ik}}$ and that each \mathbf{z}_i is independent and identically distributed according to a multinomial distribution of one draw on G categories with probabilities τ_1, \dots, τ_G . The resulting *complete-data loglikelihood* is

$$l(\theta_k, \tau_k, \mathbf{z}_{ik} | \mathbf{x}) = \sum_{i=1}^n \sum_{k=1}^G z_{ik} [\log \tau_k f_k(\mathbf{x}_i | \theta_k)]. \quad (6)$$

The quantity $\hat{z}_{ik} = E[z_{ik} | \mathbf{x}_i, \theta_1, \dots, \theta_G]$ for model (6) is the conditional expectation of z_{ik} given the observation \mathbf{x}_i and parameter values. The value z_{ik}^* of \hat{z}_{ik} at a maximum of (2) is the conditional probability that observation i belongs to group k ; the classification of an observation \mathbf{x}_i is taken to be $\{j | z_{ij}^* = \max_k z_{ik}^*\}$.

¹Conventions for normalizing λ_k and A_k include requiring $|A_k| = 1$ [5], so that $\lambda_k = |\Sigma_k|^{1/p}$, or requiring $\max(A_k) = 1$ [2], so that λ_k is the largest eigenvalue of Σ_k .

TABLE 1. Parametrizations of the covariance matrix Σ_k in the Gaussian model and their geometric interpretation. The models shown here are those discussed in Banfield and Raftery [2].

Σ_k	Distribution	Volume	Shape	Orientation	Reference
λI	Spherical	Equal	Equal	NA	1, 2, 5, 20
$\lambda_k I$	Spherical	Variable	Equal	NA	2, 5
$\lambda D A D$	Ellipsoidal	Equal	Equal	Equal	2, 5, 21, 22
$\lambda_k D_k A_k D_k$	Ellipsoidal	Variable	Variable	Variable	2, 5, 22
$\lambda D_k A D_k$	Ellipsoidal	Equal	Equal	Variable	1, 2, 5
$\lambda_k D_k A D_k$	Ellipsoidal	Variable	Equal	Variable	2, 5

initialize \hat{z}_{ik} (this can be from a discrete classification (5))

repeat

M-step: compute maximum-likelihood parameter estimates given \hat{z}_{ik}

$$n_k \leftarrow \sum_{i=1}^n \hat{z}_{ik}$$

$$\hat{\tau}_k \leftarrow \frac{n_k}{n}$$

$$\hat{\mu}_k \leftarrow \frac{\sum_{i=1}^n \hat{z}_{ik} \mathbf{x}_i}{n_k}$$

$$\hat{\Sigma}_k : \text{depends on the model—see Celeux and Govaert [5]}$$

E-step: compute \hat{z}_{ik} given the parameter estimates from the M-step

$$\hat{z}_{ik} \leftarrow \frac{\hat{\tau}_k f_k(\mathbf{x}_i \mid \hat{\mu}_k, \hat{\Sigma}_k)}{\sum_{j=1}^G \hat{\tau}_j f_j(\mathbf{x}_i \mid \hat{\mu}_j, \hat{\Sigma}_j)}, \text{ where } f_k \text{ has the form (3)}$$

until convergence criteria are satisfied

FIGURE 2. EM algorithm for clustering via Gaussian mixture models. The strategy described in this paper initializes the iteration with indicator variables (5) corresponding to partitions from hierarchical clustering and terminates when the relative difference between successive values of the mixture loglikelihood falls below a small threshold.

The EM algorithm iterates between an E-step in which values of \hat{z}_{ik} are computed from the data with the current parameter estimates and an M-step in which the complete-data loglikelihood (6), with each z_{ik} replaced by its current conditional expectation \hat{z}_{ik} , is maximized with respect to the parameters (see Figure 2). Celeux and Govaert [5] detail both the E- and M-steps for the case of multivariate normal mixture models parametrized via the eigenvalue decomposition in (4). Under certain conditions (McLachlan and Krishnan [24], Boyles [25], Wu [26]), the method can be shown to converge to a local maximum of the mixture likelihood (2). Although the conditions under which convergence has been proven do not always hold in practice, the method is widely used in the mixture modelling context with good results. Moreover, for each observation i , $(1 - \max_k \hat{z}_{ik}^*)$ is a measure of uncertainty in the associated classification (Bensmail *et al.* [27]).

The EM algorithm for clustering has a number of limitations. First, the rate of convergence can be very slow.

This does not appear to be a problem in practice for well-separated mixtures when started with reasonable values. Second, the number of conditional probabilities associated with each observation is equal to the number of components in the mixture, so that the EM algorithm for clustering may not be practical for models with very large numbers of components. Finally, EM breaks down when the covariance matrix corresponding to one or more components becomes ill-conditioned (singular or nearly singular). In general it cannot proceed if clusters contain only a few observations or if the observations they contain are very nearly collinear. If EM for a model having a certain number of components is applied to a mixture in which there are actually fewer groups, then it may fail due to ill-conditioning.

A number of variants of the EM algorithm for clustering presented here have been studied. These include the *stochastic EM* or SEM algorithm (Broniatowski, Celeux and Diebolt [28], Celeux and Diebolt [29]), in which \hat{z}_{ik} are simulated rather than estimated in the E-step, and the

classification EM or CEM algorithm (Celeux and Govaert [30]), which converts \hat{z}_{ik} from the E-step to a discrete classification before performing the M-step. The standard k -means algorithm can be shown to be a version of the CEM algorithm corresponding to the uniform spherical Gaussian model $\Sigma_k = \lambda I$ [30].

2.4. Bayesian model selection in clustering

One advantage of the mixture-model approach to clustering is that it allows the use of approximate Bayes factors to compare models. This gives a systematic means of selecting not only the parametrization of the model (and hence the clustering method), but also the number of clusters. For a recent review of Bayes factors emphasizing the underlying concepts and scientific applications, see Kass and Raftery [31].

The Bayes factor is the posterior odds for one model against the other assuming neither is favoured *a priori*. Banfield and Raftery [2] used a heuristically derived approximation to twice the log Bayes factor, called the 'AWE', to determine the number of clusters in hierarchical clustering based on the classification likelihood. When EM is used to find the maximum mixture likelihood, a more reliable approximation to twice the log Bayes factor called the BIC (Schwarz [32]) is applicable:

$$\begin{aligned} 2 \log p(x|\mathcal{M}) + \text{constant} \\ \approx 2l_{\mathcal{M}}(x, \hat{\theta}) - m_{\mathcal{M}} \log(n) \equiv \text{BIC}, \end{aligned}$$

where $p(x|\mathcal{M})$ is the (integrated) likelihood of the data for the model \mathcal{M} , $l_{\mathcal{M}}(x, \hat{\theta})$ is the maximized mixture loglikelihood for the model and $m_{\mathcal{M}}$ is the number of independent parameters to be estimated in the model. The number of clusters is not considered an independent parameter for the purposes of computing the BIC. If each model is equally likely *a priori*, then $p(x|\mathcal{M})$ is proportional to the posterior probability that the data conform to the model \mathcal{M} . Accordingly, the larger the value of the BIC, the stronger the evidence for the model.²

The fit of a mixture model to a given data set can only improve (and the likelihood can only increase) as more terms are added to the model. Hence likelihood cannot be used directly in assessment of models for cluster analysis. In the BIC, a term is added to the loglikelihood penalizing the complexity of the model, so that it may be maximized for more parsimonious parametrizations and smaller numbers of groups than the loglikelihood. The BIC can be used to compare models with differing parametrizations, differing numbers of components, or both. Bayesian criteria other than the BIC have been used in cluster analysis (e.g. Bock [17], Binder [33]). Although regularity conditions for the BIC do not hold for mixture models, there is considerable

theoretical and practical support for its use in this context [3, 4, 6, 34, 35].

A standard convention for calibrating BIC differences is that differences of less than two correspond to weak evidence, differences between two and six to positive evidence, differences between six and ten to strong evidence, and differences greater than ten to very strong evidence (Kass and Raftery [31], Jeffreys [36]).

2.5. Model-based strategy for clustering

In practice, agglomerative hierarchical clustering based on the classification likelihood (1) with Gaussian terms often gives good, but suboptimal partitions. The EM algorithm can refine partitions when started sufficiently close to the optimal value. Dasgupta and Raftery [3] were able to obtain good results in a number of examples by using the partitions produced by model-based hierarchical agglomeration as starting values for an EM algorithm for constant-shape Gaussian models, together with the BIC to determine the number of clusters. Their approach forms the basis for a more general model-based strategy for clustering.

- Determine a maximum number of clusters to consider (M) and a set of candidate parametrizations of the Gaussian model to consider. In general M should be as small as possible.
- Do agglomerative hierarchical clustering for the unconstrained Gaussian model,³ and obtain the corresponding classifications for up to M groups.
- Do EM for each parametrization and each number of clusters $2, \dots, M$, starting with the classification from hierarchical clustering.
- Compute the BIC for the one-cluster model for each parametrization and for the mixture likelihood with the optimal parameters from EM for $2, \dots, M$ clusters. This gives a matrix of BIC values corresponding to each possible combination of parametrization and number of clusters.
- Plot the BIC values for each model. A decisive first local maximum indicates strong evidence for a model (parametrization + number of clusters).

It is important to avoid applying this procedure to a larger number of components than necessary. One reason for this is to minimize computational effort; other reasons have been discussed in Section 2.3. A heuristic that works well in practice is to select the number of clusters corresponding to the first decisive local maximum (if any) over all the parametrizations considered. There may in some cases be local maxima giving larger values of BIC due to ill-conditioning rather than a genuine indication of a better model (for further discussion, see Section 3.1).

²Kass and Raftery [31] and other authors define the BIC to have opposite sign to that given here, in which case the smaller (more negative) the BIC, the stronger the evidence for the model. We have chosen to reverse this convention in order to make it easier to interpret the plots of BIC values that we present later.

³While there is a hierarchical clustering method corresponding to each parametrization of the Gaussian model, it appears to be sufficient in practice to use only the unconstrained model for initialization.

2.6. Modelling noise and outliers

Although the model-based strategy for cluster analysis as described in Section 2.5 is not directly applicable to noisy data, the model can be modified so that EM works well with a reasonably good initial identification of the noise and clusters. Noise is modelled as a constant-rate Poisson process, resulting in the mixture likelihood

$$\begin{aligned} \tilde{\mathcal{L}}_M(\theta_1, \dots, \theta_G; \tau_0, \tau_1, \dots, \tau_G | \mathbf{x}) \\ = \prod_{i=1}^n \left[\frac{\tau_0}{V} + \sum_{k=1}^G \tau_k f_k(\mathbf{x}_i | \theta_k) \right], \end{aligned} \quad (7)$$

where V is the hypervolume of the data region, $\sum_{k=0}^G \tau_k = 1$ and each $f_k(\mathbf{x}_i | \theta_k)$ is multivariate normal. An observation contributes $1/V$ if it belongs to the noise; otherwise it contributes a Gaussian term.

The basic model-based procedure for noisy data is as follows. First, it is necessary to obtain an initial estimate of the noise. Possible approaches to de-noising include the nearest-neighbour method of Byers and Raftery [37] and the method of Allard and Fraley [38], which uses Voronoi tessellations. Next, hierarchical clustering is applied to the de-noised data. In a final step, EM based on the augmented model (7) is applied to the entire data set with the Gaussian components initialized with the hierarchical clustering partitions and the noise component initialized with the result of the de-noising procedure. The BIC is then used to select the best model representing the data.

3. EXAMPLES

3.1. Diabetes diagnosis

In this section we illustrate the model-based approach to clustering using a three-dimensional data set involving 145 observations used for diabetes diagnosis (Reaven and Miller [8]). Figure 3 is a pairs plot showing the clinical classification, which partitions the data into three groups. The variables have the following meanings:

glucose—plasma glucose response to oral glucose,
insulin—plasma insulin response to oral glucose,
sspg—degree of insulin resistance.

The clusters are overlapping and are far from spherical in shape. As a result, many clustering procedures would not work well for this application. For example, Figure 3 shows the (1, 3) projection of three-cluster classifications obtained by the single-link (nearest-neighbour) method, standard k -means and the model-based method for an unconstrained Gaussian mixture. Of the possible group assignments, those shown were chosen so as to minimize the error rate in each case. The assumption of three classes is artificial for single link and k -means, while for the model-based method the BIC was used to determine the number of groups (see later).

Neither standard k -means nor single link perform well in this example. Two of the clusters identified by single link are singletons, so that nearly all the data are assigned to one class. While all three classes resulting from standard

k -means are non-trivial, two classes are confined to one of the long thin extensions while the third class subsumes the other extension as well as their conjunction. In the clinical classification, each of the two long extensions roughly represents a cluster, while the third cluster is concentrated closer to the origin. Most clustering methods that are currently in common use work well when clusters are well separated, but many break down when clusters overlap or intersect.

It is important, however, to distinguish between single-link clustering and nearest-neighbour *discrimination*. In discrimination, there is a ‘training set’ of data whose group memberships are known in advance, while in clustering, all group memberships are unknown. Nearest-neighbour discrimination assigns a data point to the same group as the point in the training set nearest to it. It often works very well (e.g. Ripley [39]), but its success depends entirely on the available training set.

Figure 4 gives a plot of the BIC for six model-based methods (spherical models with equal and varying volumes, constant variance, unconstrained variance and constant shape models with equal and varying volumes). The first local maximum (in this case also the global maximum) occurs for the unconstrained model with three clusters, for which the classification assignment is shown in Figure 1. For initial values in EM, we used z_{ik} given by equation (5) for the discrete classification from agglomerative hierarchical clustering for the unconstrained model ($\lambda_k D_k A_k D_k^T$) in all cases, leaving the model selection to the EM phase.

Of note is that no values of the BIC are given in Figure 4 for the spherical, varying-volume model for nine clusters and for the unconstrained model for eight and nine clusters. In these cases, the covariance matrix associated with one or more of the mixture components is ill-conditioned, so that the loglikelihood and hence the BIC cannot be computed. Hierarchical clustering for the spherical, varying-volume model produces a nine-cluster solution in which one cluster is a singleton, and for the unconstrained model it produces eight- and nine-cluster solutions in which one cluster contains three points. Because the data are three-dimensional, a minimum of four points is required for the estimate of the covariance matrix to be non-singular. The algorithms used for EM and for computing the BIC monitor an estimate of the reciprocal condition number (smallest to largest eigenvalue ratio) of the covariances. This latter quantity falls in the range $[0, 1]$ and values near zero imply ill-conditioning [40]. Computations are less reliable for ill-conditioned problems and as a result ill-conditioning may cause anomalies before reaching the point of actual failure. In our implementation, EM terminates with a warning if one or more estimated covariance matrices are judged to be too close to singularity, and the BIC calculation produces a missing value under the same circumstances. Table 2 shows reciprocal condition estimates for six different Gaussian mixture models for up to nine clusters. It should also be clear that EM started from partitions obtained by hierarchical clustering should not be continued for higher numbers of clusters once ill-conditioning is encountered.

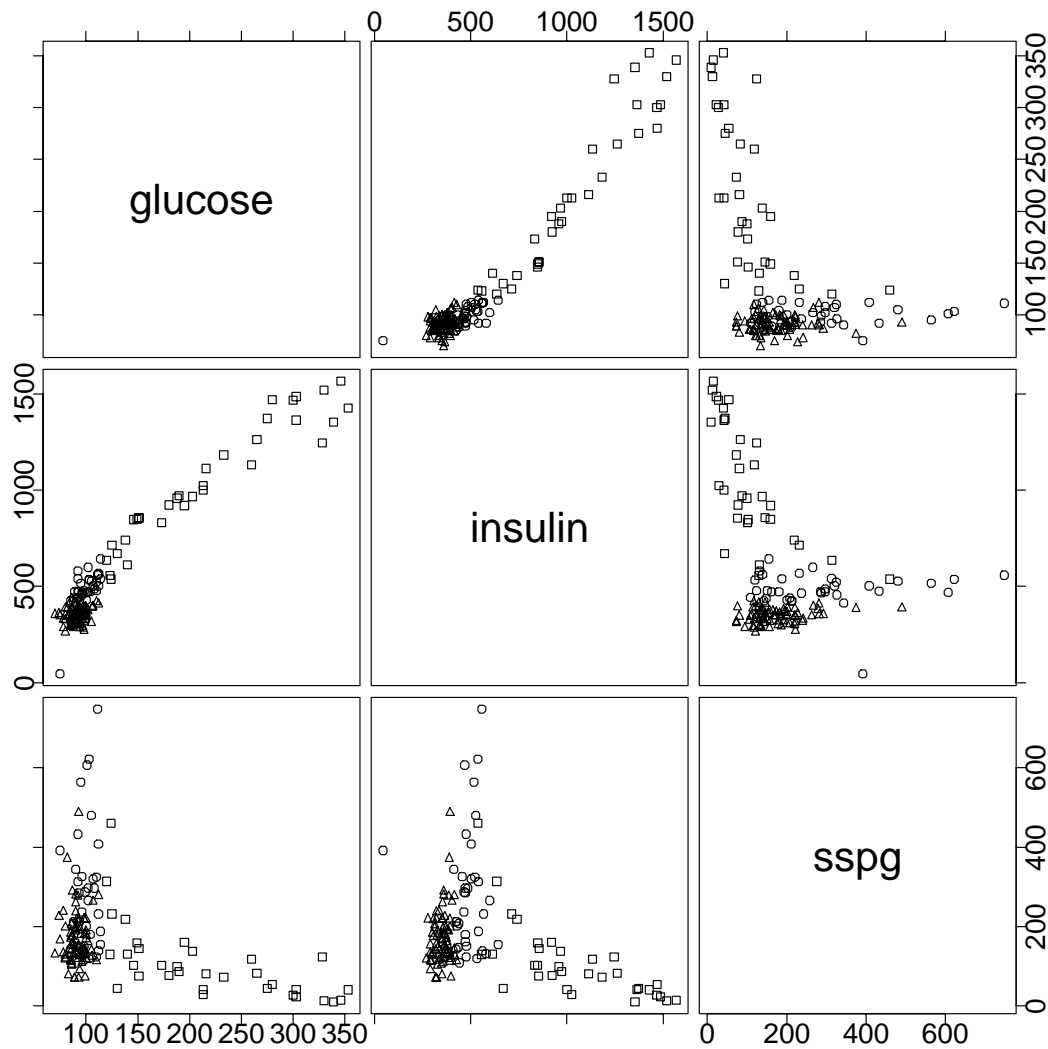


FIGURE 3. Pairs plot showing the clinical classification of the diabetes data. The symbols have the following interpretation: triangles—normal; circles—chemical diabetes; squares—overt diabetes.

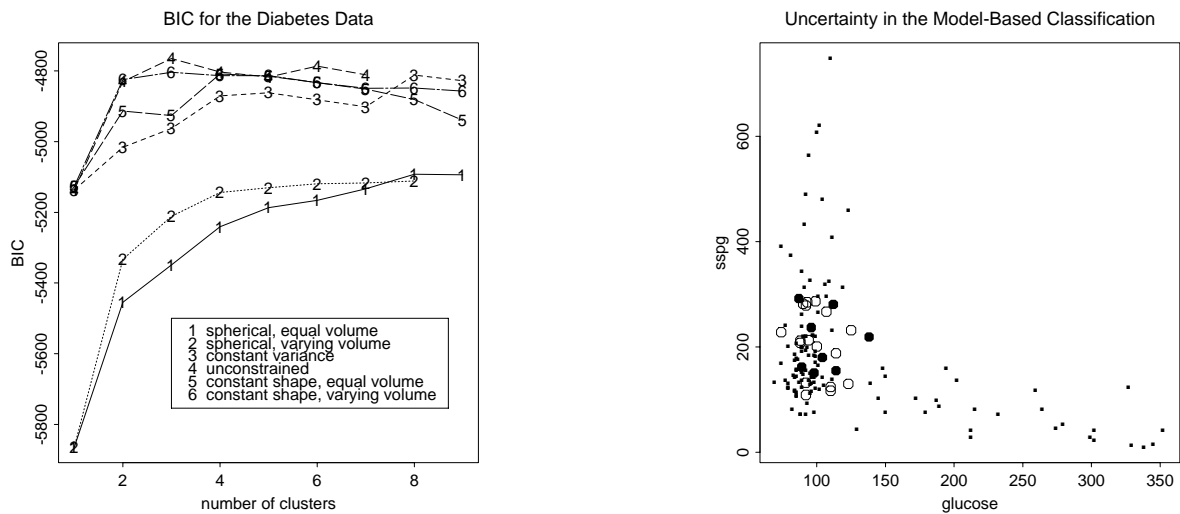


FIGURE 4. The plot on the left shows the Bayesian information criterion (BIC) for model-based methods applied to the diabetes data. The first local (also global) maximum occurs for the unconstrained model with three clusters. The plot on the right depicts the uncertainty of the classification produced by the best model (unconstrained, three clusters) indicated by the BIC. The symbols have the following interpretation: dots <0.1 ; open circles ≥ 0.1 and <0.2 ; filled circles ≥ 0.2 .

TABLE 2. Minimum reciprocal condition estimates for covariances in model-based methods applied to the diabetes data. Rows correspond to models and columns to numbers of components. Values near zero are cases in which there is either a very small cluster or one whose points are very nearly collinear.

	1	2	3	4	5	6	7	8	9
λI	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0
$\lambda_k I$	1.0	1.0	1.0	1.0	1.0	1.0	1.0	1.0	0.0
Σ	0.33	0.062	0.10	0.047	0.053	0.053	0.053	0.027	0.031
Σ_k	0.33	0.020	0.0060	0.0064	0.0067	10^{-7}	10^{-7}	10^{-32}	10^{-32}
$\lambda D_k A D_k^T$	0.0025	0.0048	0.0044	0.0049	0.0072	0.0070	0.0070	0.0017	0.0024
$\lambda_k D_k A D_k^T$	0.0025	0.0035	0.0070	0.0065	0.0065	0.0063	0.0046	0.0039	0.0027

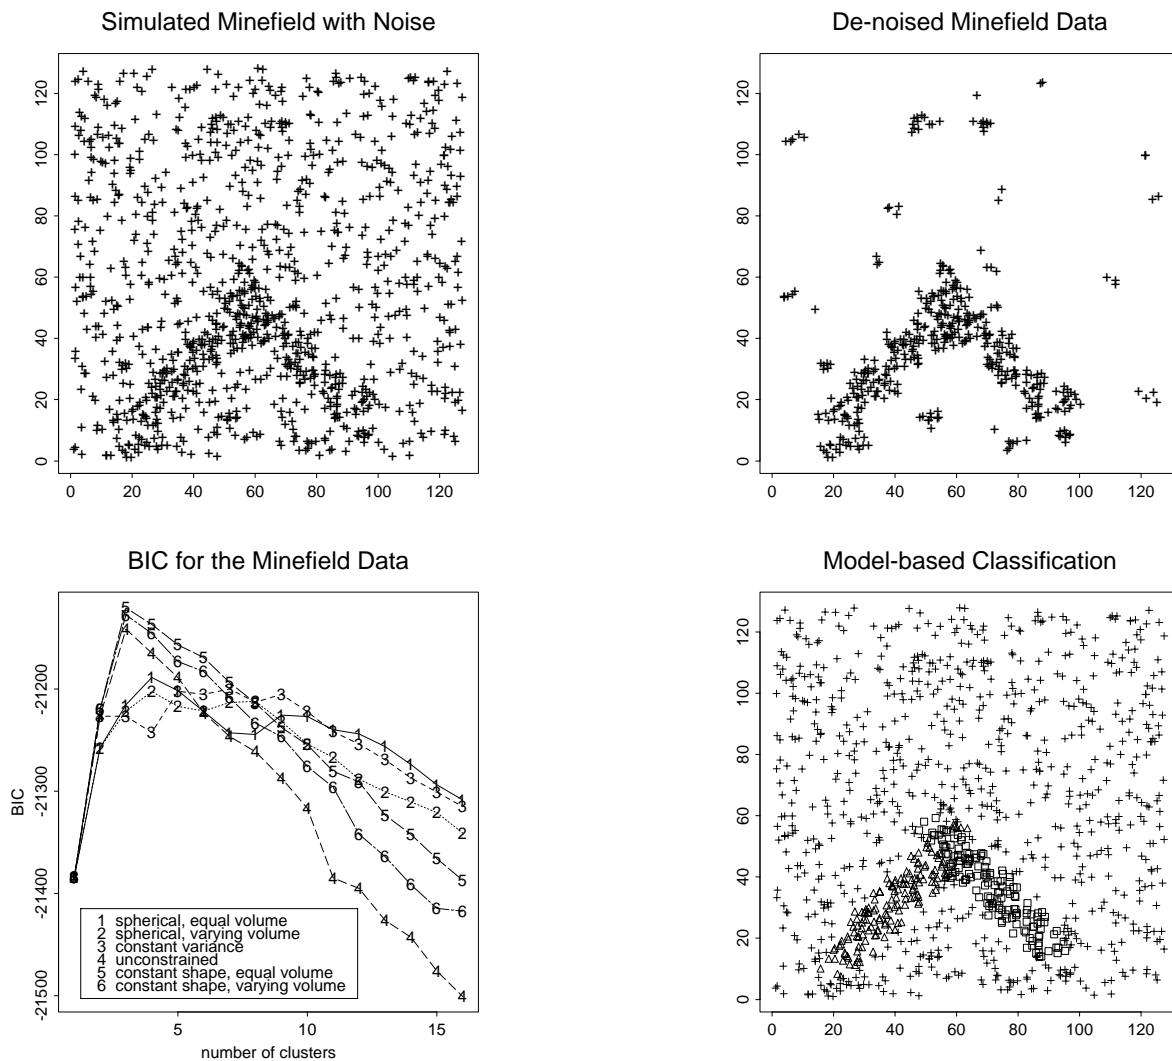


FIGURE 5. Model-based classification of a simulated minefield with noise. Hierarchical clustering was first applied to data after five nearest neighbour de-noising. EM was then applied to the full data set with the noise term included in the model.

3.2. Minefield detection in the presence of noise

Figure 5 shows the results of the model-based strategy for noise (Section 2.6) on simulated minefield data (Muise and Smith [41], see also [3]). The data arise from the processing of a series of images taken by a reconnaissance aircraft in which a large number of points are identified as representing

possible mines, but many of these are in fact false positives (noise). The assumption is that the imaged area does not lie completely within a minefield and that if there is a minefield it will occur in an area where there is a higher density of identified points. The goals are to determine whether the image contains one or more minefields and to give the location of any minefields that may be present.

The initial de-noising for Figure 5 was carried out using the `NNclean` procedure for nearest-neighbour de-noising [37]. The BIC is clearly maximized at a value of three (two clusters plus noise) and favours the uniform-shape, equal-volume model. The two clusters together give an accurate reconstruction of the actual minefield.

It should be noted that the method is sensitive to the value of V , the assumed volume of the data region. Here it is clear that V is the area of the image; Banfield and Raftery [2] and Dasgupta and Raftery [3] similarly used the volume of the smallest hyperrectangle with sides parallel to the axes that contains all the data points. However, this value could overestimate V in many cases. Another possibility is to take V to be the smallest hyperrectangle with sides parallel to the principal components of the data that contains all the data points. Our implementation uses the smaller of these two alternatives as a default, but also allows specification of V by the user. A better solution might be to use the volume of the convex hull of the data, although this may not be practical to compute in higher dimensions.

4. SOFTWARE

Software implementing state-of-the-art algorithms for hierarchical clustering [23] and EM based on the various parametrizations of Gaussian clustering models is available through the internet—for details see http://www.stat.washington.edu/fraley/mclust_soft.shtml. Included are functions that incorporate hierarchical clustering, EM and BIC in the model-based cluster analysis strategy described in this paper. This software is designed to interface with the commercial interactive software S-PLUS.⁴ An earlier version of the model-based hierarchical clustering software is included in the S-PLUS package as the function `mclust`. Subscription information for a mailing list for occasional announcements such as software updates can also be found on the same web page.

An S-PLUS function `NNclean` implementing the nearest-neighbour de-noising method [37] is available at <http://lib.stat.cmu.edu/S/nnclean>.

5. DISCUSSION

We have described a clustering methodology based on multivariate normal mixture models and shown that it can give much better performance than existing methods. This approach uses model-based agglomerative hierarchical clustering to initialize the EM algorithm for a variety of models and applies Bayesian model selection methods to determine the best clustering method along with the number of clusters. The uncertainty associated with the final classification can be assessed through the conditional probabilities from EM.

This approach has some limitations, however. The first is that computational methods for hierarchical clustering

have storage and time requirements that grow at a faster than linear rate relative to the size of the initial partition, so that they cannot be directly applied to large data sets. One way around this is to determine the structure of some subset of the data according to the strategy given here and either use the resulting parameters as initial values for EM with all of the data or classify the remaining observations via supervised classification or discriminant analysis [2]. Bensmail and Celeux [42] have developed a method for regularized discriminant analysis based on the full range of parametrizations of Gaussian mixtures (4). Alternatively, fast methods for determining an initial rough partition can be used to reduce computational requirements. Posse [43] suggested a method based on the minimum spanning tree for this purpose and has shown that it works well in practice.

Second, although experience to date suggests that models based on multivariate normal distribution are sufficiently flexible to accommodate many practical situations, the underlying assumption is that groups are concentrated locally about linear subspaces, so that other models or methods may be more suitable in some instances. In Section 3.2, we obtained good results on noisy data by combining the model-based methodology with a separate de-noising procedure. This example also suggests that nonlinear features can in some instances be well represented in the present framework as piecewise linear features, using several groups. There are alternative models in which classes are characterized by different geometries such as linear manifolds (e.g. Bock [17], Diday [44], Späth [45]). When features are strongly curvilinear, curves about which groups are centred can be modelled by using *principal curves* (Hastie and Stuetzle [46]). Clustering about principal curves has been successfully applied to automatic identification of ice-floe contours [47, 48], tracking of ice floes [49] and modelling ice-floe leads [50]. Initial estimation of ice-floe outlines is accomplished by means of mathematical morphology (e.g. [51]). Principal curve clustering in the presence of noise using BIC is discussed in Stanford and Raftery [52].

In situations where the BIC is not definitive, more computationally intensive Bayesian analysis may provide a solution. Bensmail *et al.* [27] showed that exact Bayesian inference via Gibbs sampling, with calculations of Bayes factors using the Laplace–Metropolis estimator, works well in several real and simulated examples.

Approaches to clustering based on the classification likelihood (1) are also known as *classification maximum likelihood* methods (e.g. McLachlan [53], Bryant and Williamson [54]) or *fixed-classification* methods (e.g. Bock [19, 55, 56]). There are alternatives to the classification and mixture likelihoods given in Section 2.2, such as the classification likelihood of Symons [57]:

$$\begin{aligned} \mathcal{L}_C(\theta_1, \dots, \theta_G; \tau_1, \dots, \tau_G; \gamma_1, \dots, \gamma_n \mid \mathbf{x}) \\ = \prod_{i=1}^n \tau_{\gamma_i} f_{\gamma_i}(\mathbf{x}_i \mid \theta_{\gamma_i}), \end{aligned}$$

⁴MathSoft Inc., Seattle, WA USA—<http://www.mathsoft.com/splus>

and the posterior likelihood of Anderson [58]:

$$\mathcal{L}_P(\theta_1, \dots, \theta_G; \tau_1, \dots, \tau_G; z_{11}, z_{12}, \dots, z_{nn} \mid \mathbf{x}) \\ = \prod_{k=1}^G \prod_{i=1}^n \tau_k^{z_{ik}} f(\mathbf{x}_i \mid \theta_k)^{z_{ik}}.$$

The former is the complete data likelihood for the EM algorithm when z_{ik} are restricted to be indicator variables (5), while the latter has the same form as the complete data likelihood for the EM algorithm, but includes z_{ik} as parameters to be estimated. Fuzzy clustering methods (Bezdek [59]), which are not model-based, also provide degrees of membership for observations.

The k -means algorithm has been applied not only to the classical sum-of-squares criterion but also to other model-based clustering criteria (e.g. Bock [17, 19, 56], Diday and Govaert [60], Diday [44], Späth [45], Celeux and Govaert [5]). Other model-based clustering methodologies include Cheeseman and Stutz [61, 62], implemented in the AutoClass software, and McLachlan *et al.* [12, 63, 64], implemented in the EMMIX (formerly MIXFIT) software. AutoClass handles both discrete data and continuous data, as well as data that has both discrete and continuous variables. Both AutoClass for continuous data and EMMIX rely on the EM algorithm for the multivariate normal distribution; EMMIX allows the choice of either equal, unconstrained or diagonal covariance matrices, while in Autoclass the covariances are assumed to be diagonal. As in our approach, AutoClass uses approximate Bayes factors to choose the number of clusters (see also Chickering and Heckerman [65]), although their approximation differs from the BIC. EMMIX determines the number of clusters by resampling, and has the option of modelling outliers by fitting mixtures of multivariate t -distributions (McLachlan and Peel [66]). In Autoclass, EM is initialized using random starting values, the number of trials being determined through specification of a limit on the running time. Options for initializing EM in EMMIX include the most common heuristic hierarchical clustering methods, as well as k -means, whereas we use the model-based hierarchical clustering solution as an initial value.

ACKNOWLEDGEMENTS

This work was funded by the Office of Naval Research under contract Nos N00014-96-1-0192 and N00014-96-1-0330. Thanks go to Simon Byers for providing the NNclean de-noising procedure and to the referees for suggesting a number of improvements.

REFERENCES

- [1] Murtagh, F. and Raftery, A. E. (1984) Fitting straight lines to point patterns. *Pattern Recognition*, **17**, 479–483.
- [2] Banfield, J. D. and Raftery, A. E. (1993) Model-based Gaussian and non-Gaussian clustering. *Biometrics*, **49**, 803–821.
- [3] Dasgupta, A. and Raftery, A. E. (1998) **Detecting features in spatial point processes with clutter via model-based clustering.** *J. Amer. Stat. Assoc.*, **93**, 294–302.
- [4] Campbell, J. G., Fraley, C., Murtagh, F. and Raftery, A. E. (1997) Linear flaw detection in woven textiles using model-based clustering. *Pattern Recognition Lett.*, **18**, 1539–1548.
- [5] Celeux, G. and Govaert, G. (1995) Gaussian parsimonious clustering models. *Pattern Recognition*, **28**, 781–793.
- [6] Mukerjee, S., Feigelson, E. D., Babu, G. J., Murtagh, F., Fraley, C. and Raftery, A. (1998) **Three types of gamma ray bursts.** *Astrophys. J.*, **508**, 314–327.
- [7] Dempster, A. P., Laird, N. M. and Rubin, D. B. (1977) Maximum likelihood for incomplete data via the EM algorithm. *J. R. Stat. Soc., B*, **39**, 1–38.
- [8] Reaven, G. M. and Miller, R. G. (1979) An attempt to define the nature of chemical diabetes using a multidimensional analysis. *Diabetologia*, **16**, 17–24.
- [9] Hartigan, J. A. (1975) *Clustering Algorithms*. Wiley, New York.
- [10] Gordon, A. D. (1981) *Classification: Methods for the Exploratory Analysis of Multivariate Data*. Chapman & Hall, London.
- [11] Murtagh, F. (1985) Multidimensional clustering algorithms. *CompStat Lectures*, **4**. Physica Verlag, Mostbach.
- [12] McLachlan, G. J. and Basford, K. E. (1988) *Mixture Models: Inference and Applications to Clustering*. Marcel Dekker, New York.
- [13] Kaufman, L. and Rousseeuw, P. J. (1990) *Finding Groups in Data*. Wiley, New York.
- [14] MacQueen, J. (1967) Some methods for classification and analysis of multivariate observations. In Cam, L. M. L. and Neyman, J. (eds), *Proc. 5th Berkeley Symp. on Mathematical Statistics and Probability*, vol. 1, pp. 281–297. University of California Press.
- [15] Hartigan, J. A. and Wong, M. A. (1978) Algorithm AS 136: a K -means clustering algorithm. *Appl. Stat.*, **28**, 100–108.
- [16] Engelman, L. and Hartigan, J. A. (1969) Percentage points of a test for clusters. *J. Amer. Stat. Assoc.*, **64**, 1647.
- [17] Bock, H. H. (1974) *Automatische Klassifikation (Clusteranalyse)*. Vandenhoeck & Ruprecht, Göttingen.
- [18] Bozdogan, H. (1993) Choosing the number of component clusters in the mixture model using a new informational complexity criterion of the inverse Fisher information matrix. In Opitz, O., Lausen, B. and Klar, R. (eds), *Information and Classification*, pp. 40–54. Springer-Verlag, Berlin.
- [19] Bock, H. H. (1996) Probability models and hypothesis testing in partitioning cluster analysis. In Arabie, P., Hubert, L. and DeSorte, G. (eds), *Clustering and Classification*, pp. 377–453. World Science Publishers, Singapore.
- [20] Ward, J. H. (1963) Hierarchical groupings to optimize an objective function. *J. Amer. Stat. Assoc.*, **58**, 234–244.
- [21] Friedman, H. P. and Rubin, J. (1967) On some invariant criteria for grouping data. *J. Amer. Stat. Assoc.*, **62**, 1159–1178.
- [22] Scott, A. J. and Symons, M. J. (1971) Clustering methods based on likelihood ratio criteria. *Biometrics*, **27**, 387–397.
- [23] Fraley, C. (1999) Algorithms for model-based Gaussian hierarchical clustering. *SIAM J. Sci. Comput.*, **20**, 270–281.
- [24] McLachlan, G. J. and Krishnan, T. (1997) *The EM Algorithm and Extensions*. Wiley, New York.
- [25] Boyles, R. A. (1983) On the convergence of the EM algorithm. *J. R. Stat. Soc., B*, **45**, 47–50.
- [26] Wu, C. F. J. (1983) On convergence properties of the EM algorithm. *Annals. Stat.*, **11**, 95–103.

- [27] Bensmail, H., Celeux, G., Raftery, A. E. and Robert, C. P. (1997) Inference in model-based cluster analysis. *Stat. Comput.*, **7**, 1–10.
- [28] Broniatowski, M., Celeux, G. and Diebolt, J. (1984) Reconnaissance de mélanges de densités par un algorithme d'apprentissage probabiliste. In Diday, E., Jambu, M., Lebart, L., Pagès, J.-P. and Tomassone, R. (eds), *Data Analysis and Informatics*, III, Elsevier Science, pp. 359–373.
- [29] Celeux, G. and Diebolt, J. (1985) The SEM algorithm: A probabilistic teacher algorithm derived from the EM algorithm for the mixture problem. *Comput. Stat. Quart.*, **2**, 73–82.
- [30] Celeux, G. and Govaert, G. (1992) A classification EM algorithm for clustering and two stochastic versions. *Comput. Stat. Data Anal.*, **14**, 315–332.
- [31] Kass, R. E. and Raftery, A. E. (1995) Bayes factors. *J. Amer. Stat. Assoc.*, **90**, 773–795.
- [32] Schwarz, G. (1978) Estimating the dimension of a model. *Ann. Stat.*, **6**, 461–464.
- [33] Binder, D. A. (1978) Bayesian cluster analysis. *Biometrika*, **65**, 31–38.
- [34] Leroux, M. (1992) Consistent estimation of a mixing distribution. *Annals. Stat.*, **20**, 1350–1360.
- [35] Roeder, K. and Wasserman, L. (1997) Practical Bayesian density estimation using mixtures of normals. *J. Amer. Stat. Assoc.*, **92**, 894–902.
- [36] Jeffreys, H. (1961) *Theory of Probability* (3rd edn). Clarendon, Oxford.
- [37] Byers, S. D. and Raftery, A. E. (1998) Nearest neighbor clutter removal for estimating features in spatial point processes. *J. Amer. Stat. Assoc.*, **93**, 577–584.
- [38] Allard, D. and Fraley, C. (1997) Nonparametric maximum likelihood estimation of features in spatial point processes using Voronoï tessellation. *J. Amer. Stat. Assoc.*, **92**, 1485–1493.
- [39] Ripley, B. D. (1994) Neural networks and related methods for classification. *J. R. Stat. Soc.*, B, **56**, 409–456.
- [40] Golub, G. H. and Van Loan, C. F. (1996) *Matrix Computations* (3rd edn). Johns Hopkins University Press, Baltimore, MD.
- [41] Muise, R. and Smith, C. (1991) *Nonparametric Minefield Detection and Localization*. Technical Report CSS-TM-591-91, Coastal Systems Station, Panama City, Florida.
- [42] Bensmail, H. and Celeux, G. (1996) Regularized Gaussian discriminant analysis through eigenvalue decomposition. *J. Amer. Stat. Assoc.*, **91**, 1743–1748.
- [43] Posse, C. (1998) *Hierarchical Model-based Clustering for Large Data Sets*. Technical report, University of Minnesota, School of Statistics.
- [44] Diday, E. (1979) *Optimisation en Classification Automatique*. INRIA, France.
- [45] Späth, H. (1985) *Cluster Dissection and Analysis: Theory, Fortran Programs, Examples*. Ellis Horwood, New York.
- [46] Hastie, T. and Stuetzle, W. (1989) Principal Curves. *J. Amer. Stat. Assoc.*, **84**, 502–516.
- [47] Banfield, J. D. and Raftery, A. E. (1992) Ice floe identification in satellite images using mathematical morphology and clustering about principal curves. *J. Amer. Stat. Assoc.*, **87**, 7–16.
- [48] Banfield, J. D. and Raftery, A. E. (1992) Identifying ice floes in satellite images. *Naval Research Rev.*, **43**, 2–18.
- [49] Banfield, J. D. (1991) Automated tracking of ice floes: A statistical approach. *IEEE Trans. Geoscience Remote Sensing*, **29**, 905–911.
- [50] Banfield, J. D. (1992) Skeletal modeling of ice leads. *IEEE Trans. Geoscience Remote Sensing*, **30**, 918–923.
- [51] Heijmans, H. J. A. M. (1995) Mathematical morphology: A modern approach in image processing based on algebra and geometry. *SIAM Rev.*, **37**, 1–36.
- [52] Stanford, D. and Raftery, A. E. (1997) *Principal Curve Clustering with Noise*. Technical Report 317, University of Washington, Department of Statistics.
- [53] McLachlan, G. (1982) The classification and mixture maximum likelihood approaches to cluster analysis. In Krishnaiah, P. R. and Kanal, L. N. (eds), *Handbook of Statistics*, Vol. 2, pp. 199–208. North-Holland, Amsterdam.
- [54] Bryant, P. and Williamson, A. J. (1986) Maximum likelihood and classification: a comparison of three approaches. In Gaul, W. and Schader, R. (eds), *Classification as a Tool of Research*, pp. 33–45. Elsevier Science, Amsterdam.
- [55] Bock, H. H. (1996) Probability models in partitional cluster analysis. *Comput. Stat. Data Anal.*, **23**, 5–28.
- [56] Bock, H. H. (1996) Probability models in partitional cluster analysis. In Ferligoj, A. and Kramberger, A. (eds), *Developments in Data Analysis*, pp. 3–25. FDV, Metodoloski zvezki 12, Ljubljana, Slovenia.
- [57] Symons, M. J. (1981) Clustering criteria and multivariate normal mixtures. *Biometrics*, **37**, 35–43.
- [58] Anderson, J. J. (1985) Normal mixtures and the number of clusters problem. *Comput. Stat. Quart.*, **2**, 3–14.
- [59] Bezdek, J. C. (1981) *Pattern Recognition with Fuzzy Objective Function Algorithms*. Plenum, New York.
- [60] Diday, E. and Govaert, G. (1974) Classification avec distances adaptives. *Comptes Rendues Acad. Sci. Paris*, **278**, 993.
- [61] Cheeseman, P. and Stutz, J. (1995) Bayesian classification (AutoClass): Theory and results. In Fayyad, U. et al. (eds), *Advances in Knowledge Discovery and Data Mining*, pp. 153–180. AAAI Press.
- [62] Stutz, J. and Cheeseman, P. (1995) AutoClass—a Bayesian approach to classification. In Skilling, J. and Sibisi, S. (eds), *Maximum Entropy and Bayesian Methods, Cambridge 1994*. Kluwer, Dordrecht.
- [63] McLachlan, G. J. and Peel, D. (1998) MIXFIT: An algorithm for automatic fitting and testing of normal mixtures. In *Proc. 14th Int. Conf. on Pattern Recognition*, vol. 1, pp. 553–557. IEEE Computer Society.
- [64] Peel, D. and McLachlan, G. J. (1998) *User's Guide to EMMIX—Version 1.0*. University of Queensland, Australia.
- [65] Chickering, D. M. and Heckerman, D. (1997) Efficient approximations for the marginal likelihood of Bayesian networks with hidden variables. *Machine Learning*, **29**, 181–244.
- [66] McLachlan, G. J. and Peel, D. (1998) Robust cluster analysis via mixtures of multivariate *t*-distributions. In Amin, A. et al. (eds), *Lecture Notes in Computer Science*, vol. 1451, pp. 658–666. Springer, Berlin.